

▼ Keras -- MLPs on MNIST (10 models trained)

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

↳ Using TensorFlow backend.

```
%matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu

↳ Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>
11493376/11490434 [=====] - 2s 0us/step

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

↳ Number of training examples : 60000 and each image is of shape (28, 28)
Number of training examples : 10000 and each image is of shape (28, 28)

```
# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

after converting the input images from 3d to 2d vectors

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

↳ Number of training examples : 60000 and each image is of shape (784)
Number of training examples : 10000 and each image is of shape (784)

```
# An example data point
print(X_train[0])
```

↳

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175 26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  0  30 36 94 154
170 253 253 253 253 253 225 172 253 242 195 64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251 93 82
 82 56 39  0  0  0  0  0  0  0  0  0  0  0  0 18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205 11  0 43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253 90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0 139 253 190 2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0 11 190 253 70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0 35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119 25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0 45 186 253 253 150 27  0  0  0  0  0  0  0  0  0  0
```

This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu

```
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0 46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0 39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0 24 114 221 253 253 253
253 201 78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0 23 66 213 253 253 253 253 198 81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 18 171 219 253 253 253 253 195
80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
55 172 226 253 253 253 253 244 133 11  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 136 253 253 253 212 135 132 16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0]
```

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
# X => (X - Xmin)/(Xmax-Xmin) = X/255
```

```
X_train = X_train/255
X_test = X_test/255
```

```
# example data point after normlizing
print(X_train[0])
```



```
[0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.]
```

This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu

```
0.      0.      0.      0.      0.      0.
0.      0.      0.11764706 0.14117647 0.36862745 0.60392157
0.66666667 0.99215686 0.99215686 0.99215686 0.99215686 0.99215686
0.88235294 0.6745098 0.99215686 0.94901961 0.76470588 0.25098039
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.19215686
0.93333333 0.99215686 0.99215686 0.99215686 0.99215686 0.99215686
0.99215686 0.99215686 0.99215686 0.98431373 0.36470588 0.32156863
0.32156863 0.21960784 0.15294118 0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.07058824 0.85882353 0.99215686
0.99215686 0.99215686 0.99215686 0.99215686 0.77647059 0.71372549
0.96862745 0.94509804 0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.31372549 0.61176471 0.41960784 0.99215686
0.99215686 0.80392157 0.04313725 0.      0.16862745 0.60392157
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.05490196 0.00392157 0.60392157 0.99215686 0.35294118
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.54509804 0.99215686 0.74509804 0.00784314 0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.04313725
0.74509804 0.99215686 0.2745098 0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.]
```

0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.1372549	0.94509804
0.88235294	0.62745098	0.42352941	0.00392157	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.31764706	0.94117647	0.99215686
0.99215686	0.46666667	0.09803922	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.17647059	0.72941176	0.99215686	0.99215686
0.58823529	0.10588235	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.0627451	0.36470588	0.98823529	0.99215686	0.73333333
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.97647059	0.99215686	0.97647059	0.25098039	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.18039216	0.50980392	0.71764706	0.99215686

This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu

[illegible]

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])
```

Softmax classifier

```
# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor:

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_regularizer=
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation:  $\text{output} = \text{activation}(\text{dot}(\text{input}, \text{kernel}) + \text{bias})$ , where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

#  $\text{output} = \text{activation}(\text{dot}(\text{input}, \text{kernel}) + \text{bias}) \Rightarrow y = \text{activation}(\text{WT} \cdot X + b)$ 

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument

# from keras.layers import Activation, Dense
```

```
# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation
```

```
# some model parameters
output_dim = 10
input_dim = X_train.shape[1]
batch_size = 128
nb_epoch = 20
```

* 2 Hidden layers MLP + ReLU + ADAM without batch normalization and d

```
model1_1 = Sequential()
model1_1.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNor
model1_1.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125
```

This file was updated remotely or in another tab. To force a save, overwriting the last update, select
Save from the File menu

```
history = model1_1.fit(X_train, y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, valid
```



Model: "sequential_8"

Layer (type)	Output Shape	Param #
dense_28 (Dense)	(None, 512)	401920
dense_29 (Dense)	(None, 128)	65664
dense_30 (Dense)	(None, 10)	1290

Total params: 468,874

Trainable params: 468,874

Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 78us/step - loss: 0.2223 - acc: 0.9

Epoch 2/20

60000/60000 [=====] - 3s 57us/step - loss: 0.0851 - acc: 0.9

Epoch 3/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0511 - acc: 0.9

Epoch 4/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0368 - acc: 0.9

Epoch 5/20

This file was updated remotely or in another tab. To force a save, overwriting the last update, select
Save from the File menu

Epoch 7/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0168 - acc: 0.9

Epoch 8/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0153 - acc: 0.9

Epoch 9/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0152 - acc: 0.9

Epoch 10/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0115 - acc: 0.9

Epoch 11/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0105 - acc: 0.9

Epoch 12/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0104 - acc: 0.9

Epoch 13/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0108 - acc: 0.9

Epoch 14/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0102 - acc: 0.9

Epoch 15/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0056 - acc: 0.9

Epoch 16/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0085 - acc: 0.9

Epoch 17/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0098 - acc: 0.9

Epoch 18/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0054 - acc: 0.9

Epoch 19/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0090 - acc: 0.9

Epoch 20/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0081 - acc: 0.9

score = model1_1.evaluate(X_test, Y_test, verbose=0)

```

print('Test score:', score[0])
print('Test accuracy:', score[1])
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# list of epoch numbers
x = list(range(1,nb_epoch+1))
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, v
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

☞ Test score: 0.09645916890637461
Test accuracy: 0.9819

```

score = model1_1.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

```

This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu

```

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, v

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

☞ Test score: 0.09645916890637461
Test accuracy: 0.9819

```

%matplotlib inline
w_after = model1_1.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')

```



```
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu

```
from keras.layers import Dense, Dropout, Flatten, Activation, BatchNormalization
model1 = Sequential()
model1.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.125)))
model1.add(BatchNormalization())
model1.add(Dropout(0.5))
model1.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125)))
model1.add(BatchNormalization())
model1.add(Dropout(0.5))
model1.add(Dense(output_dim, activation='softmax'))
print(model1.summary())
model1.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model1.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```



Model: "sequential_9"

Layer (type)	Output Shape	Param #
dense_31 (Dense)	(None, 512)	401920
batch_normalization_14 (Batch Normalization)	(None, 512)	2048
dropout_11 (Dropout)	(None, 512)	0
dense_32 (Dense)	(None, 128)	65664
batch_normalization_15 (Batch Normalization)	(None, 128)	512
dropout_12 (Dropout)	(None, 128)	0
dense_33 (Dense)	(None, 10)	1290
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu

Epoch 3/20

60000/60000 [=====] - 5s 91us/step - loss: 0.1773 - acc: 0.9

Epoch 4/20

60000/60000 [=====] - 5s 90us/step - loss: 0.1522 - acc: 0.9

Epoch 5/20

60000/60000 [=====] - 6s 93us/step - loss: 0.1322 - acc: 0.9

Epoch 6/20

60000/60000 [=====] - 6s 94us/step - loss: 0.1156 - acc: 0.9

Epoch 7/20

60000/60000 [=====] - 6s 93us/step - loss: 0.1066 - acc: 0.9

Epoch 8/20

60000/60000 [=====] - 5s 90us/step - loss: 0.0977 - acc: 0.9

Epoch 9/20

60000/60000 [=====] - 5s 89us/step - loss: 0.0927 - acc: 0.9

Epoch 10/20

60000/60000 [=====] - 6s 93us/step - loss: 0.0870 - acc: 0.9

Epoch 11/20

60000/60000 [=====] - 5s 88us/step - loss: 0.0819 - acc: 0.9

Epoch 12/20

60000/60000 [=====] - 5s 89us/step - loss: 0.0813 - acc: 0.9

Epoch 13/20

60000/60000 [=====] - 5s 90us/step - loss: 0.0746 - acc: 0.9

Epoch 14/20

60000/60000 [=====] - 5s 91us/step - loss: 0.0713 - acc: 0.9

Epoch 15/20

60000/60000 [=====] - 6s 93us/step - loss: 0.0659 - acc: 0.9

Epoch 16/20

60000/60000 [=====] - 6s 92us/step - loss: 0.0647 - acc: 0.9

Epoch 17/20

60000/60000 [=====] - 6s 93us/step - loss: 0.0616 - acc: 0.9

Epoch 18/20


60000/60000 [=====] - 5s 91us/step - loss: 0.0570 - acc: 0.9

Epoch 19/20


```
60000/60000 [=====] - 5s 91us/step - loss: 0.0564 - acc: 0.9
Epoch 20/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0543 - acc: 0.9
```

```
score = model1.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# list of epoch numbers
x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.052021935435730846
Test accuracy: 0.9836




This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu



```
%matplotlib inline
w_after = model1.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w, color='b')
plt.xlabel('Hidden Layer 1')
plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w, color='y')
plt.xlabel('Output Layer ')
plt.show()
```

↗



* 2 Hidden layers MLP + SIGMOID + SGD + Batch Normalization + Dropout

```
from keras.layers import Dense, Dropout, Flatten, Activation, BatchNormalization
model1 = Sequential()
model1.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(0.01, 0.01)))
model1.add(BatchNormalization())
```

This file was updated remotely or in another tab. To force a save, overwriting the last update, select
Save from the File menu

12

```
model1.add(Dropout(0.4))
model1.add(Dense(output_dim, activation='softmax'))
print(model1.summary())
model1.compile(optimizer='SGD', loss='categorical_crossentropy', metrics=['accuracy'])
history = model1.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```



WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 512)	401920
batch_normalization_1 (Batch Normalization)	(None, 512)	2048
dropout_1 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 128)	65664

This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu

dense_3 (Dense)	(None, 10)	1290
-----------------	------------	------

=====
 Total params: 471,434
 Trainable params: 470,154
 Non-trainable params: 1,280

None

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/optimizers.py:79

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 9s 149us/step - loss: 1.0443 - acc: 0.

Epoch 2/20

60000/60000 [=====] - 4s 72us/step - loss: 0.5988 - acc: 0.8

Epoch 3/20

60000/60000 [=====] - 4s 72us/step - loss: 0.5188 - acc: 0.8

Epoch 4/20

60000/60000 [=====] - 4s 74us/step - loss: 0.4803 - acc: 0.8

Epoch 5/20

60000/60000 [=====] - 4s 71us/step - loss: 0.4539 - acc: 0.8

Epoch 6/20

60000/60000 [=====] - 4s 74us/step - loss: 0.4333 - acc: 0.8

Epoch 7/20

60000/60000 [=====] - 4s 72us/step - loss: 0.4198 - acc: 0.8

Epoch 8/20

60000/60000 [=====] - 4s 74us/step - loss: 0.4062 - acc: 0.8

Epoch 9/20

60000/60000 [=====] - 4s 73us/step - loss: 0.3955 - acc: 0.8

Epoch 10/20

```

60000/60000 [=====] - 4s 73us/step - loss: 0.3863 - acc: 0.8
Epoch 11/20
60000/60000 [=====] - 4s 73us/step - loss: 0.3762 - acc: 0.8
Epoch 12/20
60000/60000 [=====] - 4s 71us/step - loss: 0.3687 - acc: 0.8
Epoch 13/20
60000/60000 [=====] - 4s 74us/step - loss: 0.3649 - acc: 0.8
Epoch 14/20
60000/60000 [=====] - 4s 72us/step - loss: 0.3617 - acc: 0.8
Epoch 15/20
60000/60000 [=====] - 4s 73us/step - loss: 0.3541 - acc: 0.8
Epoch 16/20
60000/60000 [=====] - 4s 73us/step - loss: 0.3461 - acc: 0.8
Epoch 17/20
60000/60000 [=====] - 4s 73us/step - loss: 0.3451 - acc: 0.8
Epoch 18/20
60000/60000 [=====] - 4s 72us/step - loss: 0.3389 - acc: 0.8
Epoch 19/20
60000/60000 [=====] - 4s 74us/step - loss: 0.3329 - acc: 0.9
Epoch 20/20
60000/60000 [=====] - 4s 74us/step - loss: 0.3312 - acc: 0.9

```

```

score = model1.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

```

This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu

```

x = list(range(1,nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.23021167929247022
Test accuracy: 0.9338


```

%matplotlib inline
w_after = model1.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')
plt.subplot(1,3,2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```

↗



* 2 Hidden layers MLP + tanh + SGD + Batch Normalization + Dropout

```
from keras.layers import Dense, Dropout, Flatten, Activation, BatchNormalization
model1 = Sequential()
model1.add(Dense(512, activation='tanh', input_shape=(input_dim,), kernel_initializer=RandomNormal(0.01, 0.01)))
model1.add(BatchNormalization())
model1.add(Dropout(0.4))
```

This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu

```
model1.add(Dense(output_dim, activation='softmax'))
print(model1.summary())
model1.compile(optimizer='SGD', loss='categorical_crossentropy', metrics=['accuracy'])
history = model1.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```



Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 512)	401920
batch_normalization_3 (Batch Normalization)	(None, 512)	2048
dropout_3 (Dropout)	(None, 512)	0
dense_5 (Dense)	(None, 128)	65664
batch_normalization_4 (Batch Normalization)	(None, 128)	512
dropout_4 (Dropout)	(None, 128)	0
dense_6 (Dense)	(None, 10)	1290
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu


```
Epoch 3/20
60000/60000 [=====] - 4s 71us/step - loss: 0.5301 - acc: 0.8
Epoch 4/20
60000/60000 [=====] - 4s 75us/step - loss: 0.4834 - acc: 0.8
Epoch 5/20
60000/60000 [=====] - 4s 75us/step - loss: 0.4525 - acc: 0.8
Epoch 6/20
60000/60000 [=====] - 4s 72us/step - loss: 0.4300 - acc: 0.8
Epoch 7/20
60000/60000 [=====] - 4s 73us/step - loss: 0.4120 - acc: 0.8
Epoch 8/20
60000/60000 [=====] - 4s 71us/step - loss: 0.3959 - acc: 0.8
Epoch 9/20
60000/60000 [=====] - 4s 73us/step - loss: 0.3889 - acc: 0.8
Epoch 10/20
60000/60000 [=====] - 4s 71us/step - loss: 0.3743 - acc: 0.8
Epoch 11/20
60000/60000 [=====] - 4s 72us/step - loss: 0.3671 - acc: 0.8
Epoch 12/20
60000/60000 [=====] - 4s 72us/step - loss: 0.3602 - acc: 0.8
Epoch 13/20
60000/60000 [=====] - 4s 72us/step - loss: 0.3506 - acc: 0.8
Epoch 14/20
60000/60000 [=====] - 4s 71us/step - loss: 0.3443 - acc: 0.8
Epoch 15/20
60000/60000 [=====] - 4s 72us/step - loss: 0.3377 - acc: 0.8
Epoch 16/20
60000/60000 [=====] - 4s 72us/step - loss: 0.3312 - acc: 0.9
Epoch 17/20
60000/60000 [=====] - 4s 72us/step - loss: 0.3258 - acc: 0.9
Epoch 18/20
60000/60000 [=====] - 4s 72us/step - loss: 0.3224 - acc: 0.9
Epoch 19/20
```




```
60000/60000 [=====] - 4s 73us/step - loss: 0.3180 - acc: 0.9
Epoch 20/20
60000/60000 [=====] - 4s 71us/step - loss: 0.3115 - acc: 0.9
```

```
score = model1.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# list of epoch numbers
x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.20571986180469393
Test accuracy: 0.9393




This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu



```
%matplotlib inline
w_after = model1.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')
plt.subplot(1,3,2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

Test score: 0.20571986180469393
Test accuracy: 0.9393



```
from keras.layers import Dense, Dropout, Flatten, Activation, BatchNormalization
model1 = Sequential()
model1.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.12)))
model1.add(BatchNormalization())
model1.add(Dropout(0.4))
model1.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.12)))
model1.add(BatchNormalization())
model1.add(Dropout(0.4))
model1.add(Dense(output_dim, activation='softmax'))
```

This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu

at



Model: "sequential_3"

Layer (type)	Output Shape	Param #
dense_7 (Dense)	(None, 512)	401920
batch_normalization_5 (Batch Normalization)	(None, 512)	2048
dropout_5 (Dropout)	(None, 512)	0
dense_8 (Dense)	(None, 128)	65664
batch_normalization_6 (Batch Normalization)	(None, 128)	512
dropout_6 (Dropout)	(None, 128)	0
dense_9 (Dense)	(None, 10)	1290
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20


This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu

```
Epoch 3/20
60000/60000 [=====] - 5s 82us/step - loss: 0.2510 - acc: 0.9
Epoch 4/20
60000/60000 [=====] - 5s 85us/step - loss: 0.2198 - acc: 0.9
Epoch 5/20
60000/60000 [=====] - 5s 84us/step - loss: 0.1926 - acc: 0.9
Epoch 6/20
60000/60000 [=====] - 5s 84us/step - loss: 0.1729 - acc: 0.9
Epoch 7/20
60000/60000 [=====] - 5s 82us/step - loss: 0.1572 - acc: 0.9
Epoch 8/20
60000/60000 [=====] - 5s 82us/step - loss: 0.1438 - acc: 0.9
Epoch 9/20
60000/60000 [=====] - 5s 83us/step - loss: 0.1276 - acc: 0.9
Epoch 10/20
60000/60000 [=====] - 5s 83us/step - loss: 0.1197 - acc: 0.9
Epoch 11/20
60000/60000 [=====] - 5s 84us/step - loss: 0.1088 - acc: 0.9
Epoch 12/20
60000/60000 [=====] - 5s 84us/step - loss: 0.1003 - acc: 0.9
Epoch 13/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0962 - acc: 0.9
Epoch 14/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0894 - acc: 0.9
Epoch 15/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0829 - acc: 0.9
Epoch 16/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0805 - acc: 0.9
Epoch 17/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0754 - acc: 0.9
Epoch 18/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0711 - acc: 0.9
Epoch 19/20
```


```
60000/60000 [=====] - 5s 82us/step - loss: 0.0669 - acc: 0.9
Epoch 20/20
60000/60000 [=====] - 5s 80us/step - loss: 0.0646 - acc: 0.9
```

```
score = model1.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# list of epoch numbers
x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06078589027107228
Test accuracy: 0.9818




This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu



```
%matplotlib inline
w_after = model1.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')
plt.subplot(1,3,2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

Test score: 0.06078589027107228
Test accuracy: 0.9818



2 layer + ADAM + DROPOUT + BATCH NORMALIZATION + TANH

```
from keras.layers import Dense, Dropout, Flatten, Activation, BatchNormalization
model1 = Sequential()
model1.add(Dense(512, activation='tanh', input_shape=(input_dim,), kernel_initializer=RandomNormal(0.01, 0.01)))
model1.add(BatchNormalization())
model1.add(Dropout(0.4))
```

This file was updated remotely or in another tab. To force a save, overwriting the last update, select
Save from the File menu

```
model1.add(Dense(output_dim, activation='softmax'))
print(model1.summary())
model1.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model1.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```



Model: "sequential_4"

Layer (type)	Output Shape	Param #
dense_10 (Dense)	(None, 512)	401920
batch_normalization_7 (Batch Normalization)	(None, 512)	2048
dropout_7 (Dropout)	(None, 512)	0
dense_11 (Dense)	(None, 128)	65664
batch_normalization_8 (Batch Normalization)	(None, 128)	512
dropout_8 (Dropout)	(None, 128)	0
dense_12 (Dense)	(None, 10)	1290
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu

Epoch 3/20

60000/60000 [=====] - 5s 84us/step - loss: 0.2183 - acc: 0.9

Epoch 4/20

60000/60000 [=====] - 5s 85us/step - loss: 0.1825 - acc: 0.9

Epoch 5/20

60000/60000 [=====] - 5s 83us/step - loss: 0.1561 - acc: 0.9

Epoch 6/20

60000/60000 [=====] - 5s 86us/step - loss: 0.1358 - acc: 0.9

Epoch 7/20

60000/60000 [=====] - 5s 85us/step - loss: 0.1208 - acc: 0.9

Epoch 8/20

60000/60000 [=====] - 5s 83us/step - loss: 0.1121 - acc: 0.9

Epoch 9/20

60000/60000 [=====] - 5s 85us/step - loss: 0.1008 - acc: 0.9

Epoch 10/20

60000/60000 [=====] - 5s 85us/step - loss: 0.0911 - acc: 0.9

Epoch 11/20

60000/60000 [=====] - 5s 84us/step - loss: 0.0885 - acc: 0.9

Epoch 12/20

60000/60000 [=====] - 5s 86us/step - loss: 0.0811 - acc: 0.9

Epoch 13/20

60000/60000 [=====] - 5s 86us/step - loss: 0.0761 - acc: 0.9

Epoch 14/20

60000/60000 [=====] - 5s 85us/step - loss: 0.0725 - acc: 0.9

Epoch 15/20

60000/60000 [=====] - 5s 85us/step - loss: 0.0695 - acc: 0.9

Epoch 16/20

60000/60000 [=====] - 5s 85us/step - loss: 0.0663 - acc: 0.9

Epoch 17/20

60000/60000 [=====] - 5s 89us/step - loss: 0.0613 - acc: 0.9

Epoch 18/20


60000/60000 [=====] - 5s 84us/step - loss: 0.0561 - acc: 0.9

Epoch 19/20


```
60000/60000 [=====] - 5s 84us/step - loss: 0.0563 - acc: 0.9
Epoch 20/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0527 - acc: 0.9
```

```
score = model1.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# list of epoch numbers
x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.0640238022508216
Test accuracy: 0.9822




This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu



```
%matplotlib inline
w_after = model1.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')
plt.subplot(1,3,2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

Test score: 0.0640238022508216
Test accuracy: 0.9822



3 hidden layer MLP + Batch-Norm on hidden Layers + AdamOptimizer but

Multilayer perceptron

<https://intoli.com/blog/neural-network-initialization/>
 # If we sample weights from a normal distribution $N(0, \sigma)$ we satisfy this condition with $\sigma = \sqrt{2/(n_i + n_{i+1})}$
 # $h1 \rightarrow \sigma = \sqrt{2/(n1+n1+1)} = 0.029 \rightarrow N(0, \sigma) = N(0, 0.029)$

This file was updated remotely or in another tab. To force a save, overwriting the last update, select
 Save from the File menu

```
from keras.layers.normalization import BatchNormalization
model_2 = Sequential()
model_2.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=10)))
model_2.add(BatchNormalization())
model_2.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=10)))
model_2.add(BatchNormalization())
model_2.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=10)))
model_2.add(BatchNormalization())
model_2.add(Dense(output_dim, activation='softmax'))
model_2.summary()
```



Model: "sequential_10"

Layer (type)	Output Shape	Param #
=====		
dense_34 (Dense)	(None, 512)	401920
batch_normalization_16 (Batch Normalization)	(None, 512)	2048
dense_35 (Dense)	(None, 256)	131328
batch_normalization_17 (Batch Normalization)	(None, 256)	1024
dense_36 (Dense)	(None, 128)	32896
batch_normalization_18 (Batch Normalization)	(None, 128)	512
dense_37 (Dense)	(None, 10)	1290
=====		
Total params: 571,018		
Trainable params: 569,226		
Non-trainable params: 1,792		

```
model_2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu

da



Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 9s 153us/step - loss: 0.2075 - acc: 0.
Epoch 2/20
60000/60000 [=====] - 7s 116us/step - loss: 0.0741 - acc: 0.
Epoch 3/20
60000/60000 [=====] - 7s 112us/step - loss: 0.0487 - acc: 0.
Epoch 4/20
60000/60000 [=====] - 7s 114us/step - loss: 0.0341 - acc: 0.
Epoch 5/20
60000/60000 [=====] - 7s 118us/step - loss: 0.0285 - acc: 0.
Epoch 6/20
60000/60000 [=====] - 7s 117us/step - loss: 0.0205 - acc: 0.
Epoch 7/20
60000/60000 [=====] - 7s 117us/step - loss: 0.0180 - acc: 0.
Epoch 8/20
60000/60000 [=====] - 7s 116us/step - loss: 0.0148 - acc: 0.
Epoch 9/20
60000/60000 [=====] - 7s 118us/step - loss: 0.0159 - acc: 0.
Epoch 10/20
60000/60000 [=====] - 7s 115us/step - loss: 0.0154 - acc: 0.
Epoch 11/20
60000/60000 [=====] - 7s 114us/step - loss: 0.0132 - acc: 0.
Epoch 12/20
60000/60000 [=====] - 7s 114us/step - loss: 0.0115 - acc: 0.
```

This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu

```
60000/60000 [=====] - 7s 116us/step - loss: 0.0107 - acc: 0.
Epoch 15/20
60000/60000 [=====] - 7s 114us/step - loss: 0.0080 - acc: 0.
Epoch 16/20
60000/60000 [=====] - 7s 115us/step - loss: 0.0073 - acc: 0.
Epoch 17/20
60000/60000 [=====] - 7s 113us/step - loss: 0.0090 - acc: 0.
Epoch 18/20
60000/60000 [=====] - 7s 113us/step - loss: 0.0086 - acc: 0.
Epoch 19/20
60000/60000 [=====] - 7s 113us/step - loss: 0.0073 - acc: 0.
Epoch 20/20
60000/60000 [=====] - 6s 108us/step - loss: 0.0051 - acc: 0.
```

```
score = model_2.evaluate(X_test, Y_test, verbose=0)
```

```

print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, v

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.07286385818796771
 Test accuracy: 0.9821



This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu



```

w_after = model_2.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(2, 2, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(2, 2, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(2, 2, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3')

plt.subplot(2, 2, 4)

```

```
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



3 HIDDEN LAYER AND RELU WITH DROPOUT AND BATCH NORMALIZATION

```
from keras.layers import Dropout
```

:

This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu

Nc

```
model_2_1.add(Dropout(0.5))
model_2_1.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormalizer(0.01)))
model_2_1.add(BatchNormalization())
model_2_1.add(Dropout(0.5))
model_2_1.add(Dense(128, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormalizer(0.01)))
model_2_1.add(BatchNormalization())
model_2_1.add(Dropout(0.5))
model_2_1.add(Dense(output_dim, activation='softmax'))
model_2_1.summary()
```



Model: "sequential_11"

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_38 (Dense)	(None, 512)	401920
batch_normalization_19 (Batch Normalization)	(None, 512)	2048
dropout_13 (Dropout)	(None, 512)	0
dense_39 (Dense)	(None, 256)	131328
batch_normalization_20 (Batch Normalization)	(None, 256)	1024
dropout_14 (Dropout)	(None, 256)	0
dense_40 (Dense)	(None, 128)	32896
batch_normalization_21 (Batch Normalization)	(None, 128)	512
dropout_15 (Dropout)	(None, 128)	0
dense_41 (Dense)	(None, 10)	1290
=====	=====	=====
Total params: 571,018		

This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu

```
model_2_1.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_2_1.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, vali
```



Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 9s 149us/step - loss: 0.4808 - acc: 0.

Epoch 2/20

60000/60000 [=====] - 7s 112us/step - loss: 0.2249 - acc: 0.

Epoch 3/20

60000/60000 [=====] - 7s 111us/step - loss: 0.1794 - acc: 0.

Epoch 4/20

60000/60000 [=====] - 7s 111us/step - loss: 0.1506 - acc: 0.

Epoch 5/20

60000/60000 [=====] - 7s 111us/step - loss: 0.1379 - acc: 0.

Epoch 6/20

60000/60000 [=====] - 7s 113us/step - loss: 0.1205 - acc: 0.

Epoch 7/20

60000/60000 [=====] - 7s 112us/step - loss: 0.1148 - acc: 0.

Epoch 8/20

60000/60000 [=====] - 7s 114us/step - loss: 0.1046 - acc: 0.

Epoch 9/20

60000/60000 [=====] - 7s 114us/step - loss: 0.0998 - acc: 0.

Epoch 10/20

60000/60000 [=====] - 7s 115us/step - loss: 0.0949 - acc: 0.

Epoch 11/20

60000/60000 [=====] - 7s 119us/step - loss: 0.0939 - acc: 0.

Epoch 12/20

60000/60000 [=====] - 7s 119us/step - loss: 0.0830 - acc: 0.

This file was updated remotely or in another tab. To force a save, overwriting the last update, select
Save from the File menu

60000/60000 [=====] - 7s 120us/step - loss: 0.0803 - acc: 0.

Epoch 15/20

60000/60000 [=====] - 7s 120us/step - loss: 0.0736 - acc: 0.

Epoch 16/20

60000/60000 [=====] - 7s 117us/step - loss: 0.0726 - acc: 0.

Epoch 17/20

60000/60000 [=====] - 7s 119us/step - loss: 0.0708 - acc: 0.

Epoch 18/20

60000/60000 [=====] - 7s 120us/step - loss: 0.0662 - acc: 0.

Epoch 19/20

60000/60000 [=====] - 7s 119us/step - loss: 0.0644 - acc: 0.

Epoch 20/20

60000/60000 [=====] - 7s 122us/step - loss: 0.0639 - acc: 0.

score = model_2_1.evaluate(X_test, Y_test, verbose=0)

```

print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, v

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

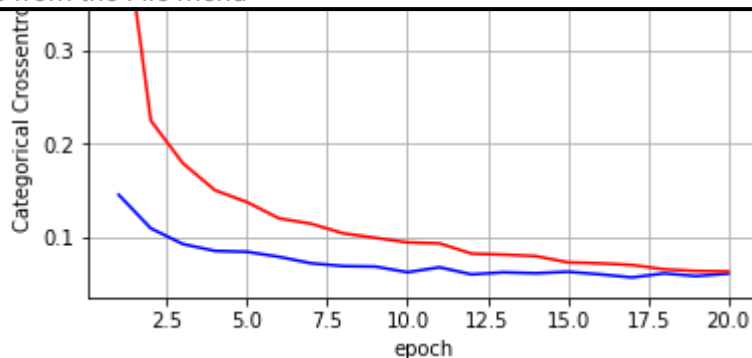
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.061630533440283033
 Test accuracy: 0.9821



This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu



```

w_after = model_2_1.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(2, 2, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(2, 2, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(2, 2, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3')

```

```
plt.subplot(2, 2, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer')
plt.show()
```



* 5 Hidden layer MLP + ReLu + AdamOptimizer

This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu

```
model_3_1.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal))
model_3_1.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal))
model_3_1.add(Dense(128, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal))
model_3_1.add(Dense(64, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal))
model_3_1.add(Dense(32, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal))
model_3_1.add(Dense(output_dim, activation='softmax'))
model_3_1.summary()
```

Model: "sequential_12"

Layer (type)	Output Shape	Param #
dense_42 (Dense)	(None, 512)	401920
dense_43 (Dense)	(None, 256)	131328
dense_44 (Dense)	(None, 128)	32896
dense_45 (Dense)	(None, 64)	8256
dense_46 (Dense)	(None, 32)	2080
dense_47 (Dense)	(None, 10)	330
Total params: 576,810		
Trainable params: 576,810		
Non-trainable params: 0		

```
model_3_1.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_3_1.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, vali
```


☞ Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 7s 108us/step - loss: 0.3922 - acc: 0.
Epoch 2/20
60000/60000 [=====] - 4s 71us/step - loss: 0.1226 - acc: 0.9
Epoch 3/20
60000/60000 [=====] - 4s 71us/step - loss: 0.0807 - acc: 0.9
Epoch 4/20
60000/60000 [=====] - 4s 70us/step - loss: 0.0589 - acc: 0.9
Epoch 5/20
60000/60000 [=====] - 4s 69us/step - loss: 0.0443 - acc: 0.9
Epoch 6/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0380 - acc: 0.9
Epoch 7/20
60000/60000 [=====] - 4s 74us/step - loss: 0.0279 - acc: 0.9
Epoch 8/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0233 - acc: 0.9
Epoch 9/20
60000/60000 [=====] - 4s 71us/step - loss: 0.0230 - acc: 0.9
Epoch 10/20
60000/60000 [=====] - 4s 71us/step - loss: 0.0209 - acc: 0.9
Epoch 11/20
60000/60000 [=====] - 4s 71us/step - loss: 0.0178 - acc: 0.9
Epoch 12/20
60000/60000 [=====] - 4s 72us/step - loss: 0.0155 - acc: 0.9
Epoch 13/20
60000/60000 [=====] - 4s 72us/step - loss: 0.0155 - acc: 0.9
Epoch 14/20
60000/60000 [=====] - 4s 72us/step - loss: 0.0172 - acc: 0.9
Epoch 15/20
60000/60000 [=====] - 4s 71us/step - loss: 0.0112 - acc: 0.9
Epoch 16/20
60000/60000 [=====] - 4s 70us/step - loss: 0.0121 - acc: 0.9
Epoch 17/20
60000/60000 [=====] - 4s 71us/step - loss: 0.0127 - acc: 0.9
Epoch 18/20
60000/60000 [=====] - 4s 70us/step - loss: 0.0127 - acc: 0.9
Epoch 19/20
60000/60000 [=====] - 4s 72us/step - loss: 0.0092 - acc: 0.9
Epoch 20/20
60000/60000 [=====] - 4s 69us/step - loss: 0.0117 - acc: 0.9
```


This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu

```


score = model_3_1.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# list of epoch numbers
x = list(range(1,nb_epoch+1))
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, v
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09843842560593985
Test accuracy: 0.9782



This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu



```

w_after = model_3_1.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(3,2 , 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(3, 2, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(3, 2, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(3,2, 4)
plt.title("Trained model Weights")

```

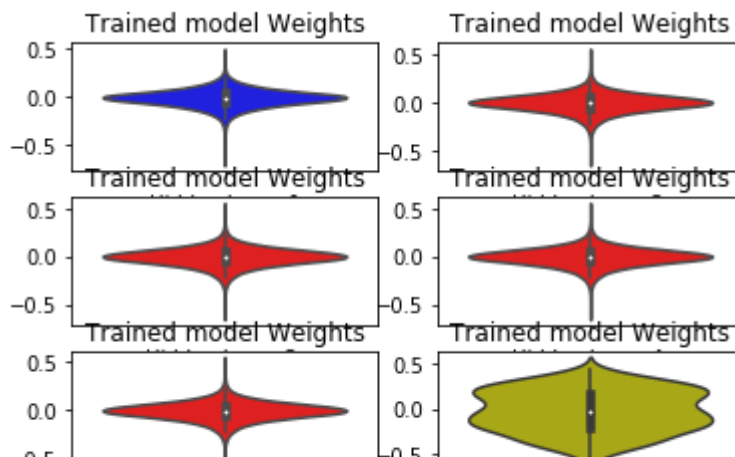
```

ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(3, 2, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(3,2, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu

* 5 Hidden layer MLP + Dropout + Batch Normalization + ReLu + AdamOpt

<https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in>

```

from keras.layers import Dropout
model_3 = Sequential()
model_3.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal))
model_3.add(BatchNormalization())
model_3.add(Dropout(0.5))
model_3.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal))
model_3.add(BatchNormalization())
model_3.add(Dropout(0.5))
model_3.add(Dense(128, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal))
model_3.add(BatchNormalization())
model_3.add(Dropout(0.5))
model_3.add(Dense(64, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal))
model_3.add(BatchNormalization())
model_3.add(Dropout(0.5))
model_3.add(Dense(32, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal))
model_3.add(BatchNormalization())
model_3.add(Dropout(0.5))
model_3.add(Dense(output_dim, activation='softmax'))
model_3.summary()

```



Model: "sequential_13"

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_48 (Dense)	(None, 512)	401920
batch_normalization_22 (Batch Normalization)	(None, 512)	2048
dropout_16 (Dropout)	(None, 512)	0
dense_49 (Dense)	(None, 256)	131328
batch_normalization_23 (Batch Normalization)	(None, 256)	1024
dropout_17 (Dropout)	(None, 256)	0
dense_50 (Dense)	(None, 128)	32896
batch_normalization_24 (Batch Normalization)	(None, 128)	512
dropout_18 (Dropout)	(None, 128)	0
dense_51 (Dense)	(None, 64)	8256
batch_normalization_25 (Batch Normalization)	(None, 64)	256

This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu

dense_52 (Dense)	(None, 32)	2080
batch_normalization_26 (Batch Normalization)	(None, 32)	128
dropout_20 (Dropout)	(None, 32)	0
dense_53 (Dense)	(None, 10)	330
=====	=====	=====
Total params: 580,778		
Trainable params: 578,794		
Non-trainable params: 1,984		

```
model_3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_val, Y_val))
```



Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 13s 225us/step - loss: 1.1314 - acc: 0
Epoch 2/20
60000/60000 [=====] - 10s 164us/step - loss: 0.4195 - acc: 0
Epoch 3/20
60000/60000 [=====] - 10s 162us/step - loss: 0.3209 - acc: 0
Epoch 4/20
60000/60000 [=====] - 10s 159us/step - loss: 0.2763 - acc: 0
Epoch 5/20
60000/60000 [=====] - 9s 158us/step - loss: 0.2495 - acc: 0.
Epoch 6/20
60000/60000 [=====] - 10s 162us/step - loss: 0.2268 - acc: 0
Epoch 7/20
60000/60000 [=====] - 10s 160us/step - loss: 0.2138 - acc: 0
Epoch 8/20
60000/60000 [=====] - 9s 156us/step - loss: 0.1929 - acc: 0.
Epoch 9/20
60000/60000 [=====] - 10s 162us/step - loss: 0.1868 - acc: 0
Epoch 10/20
60000/60000 [=====] - 10s 163us/step - loss: 0.1752 - acc: 0
Epoch 11/20
60000/60000 [=====] - 10s 163us/step - loss: 0.1716 - acc: 0
Epoch 12/20
60000/60000 [=====] - 10s 164us/step - loss: 0.1643 - acc: 0
```

This file was updated remotely or in another tab. To force a save, overwriting the last update, select
Save from the File menu

```
60000/60000 [=====] - 9s 157us/step - loss: 0.1486 - acc: 0.
Epoch 15/20
60000/60000 [=====] - 10s 158us/step - loss: 0.1457 - acc: 0
Epoch 16/20
60000/60000 [=====] - 10s 161us/step - loss: 0.1430 - acc: 0
Epoch 17/20
60000/60000 [=====] - 10s 160us/step - loss: 0.1317 - acc: 0
Epoch 18/20
60000/60000 [=====] - 10s 163us/step - loss: 0.1314 - acc: 0
Epoch 19/20
60000/60000 [=====] - 10s 163us/step - loss: 0.1263 - acc: 0
Epoch 20/20
60000/60000 [=====] - 10s 163us/step - loss: 0.1238 - acc: 0
```

```
score = model_3.evaluate(X_test, Y_test, verbose=0)
```

```

print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, v

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.07888613095330074
Test accuracy: 0.9813



This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu



```

w_after = model_3.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(3,2, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(3, 2, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(3, 2, 3)
plt.title("Trained model Weights")

```

```

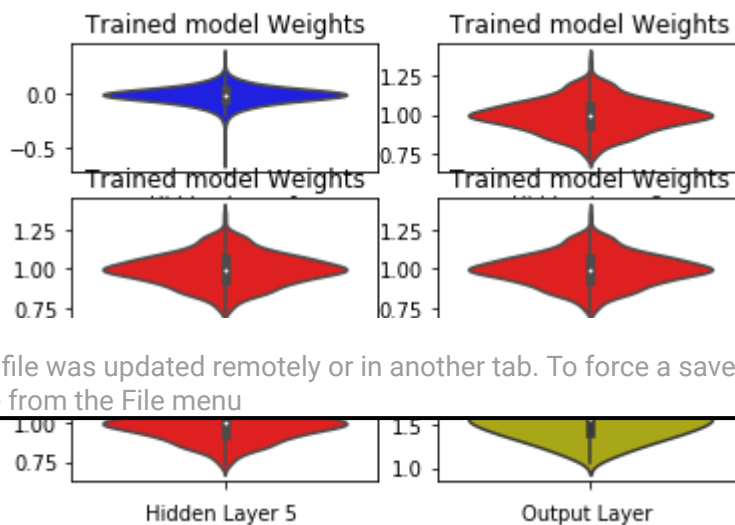
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(3,2, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(3, 2, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(3,2, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu

```

from prettytable import PrettyTable

x = PrettyTable()

x.field_names = ["Hiddden layers", "batch normalization", "activation function", "optimizer", "drop

x.add_row([2, 'no', "relu", 'adam', 'no', 98.19])
x.add_row([2, 'yes', 'relu', "adam", 'yes', 98.36])
x.add_row([2, 'yes', "sigmoid", 'SGD', 'yes', 93.38])
x.add_row([2, 'yes', 'tanh', "SGD", 'yes', 93.93])
x.add_row([2, 'yes', "sigmoid", 'adam', 'yes', 98.18])
x.add_row([2, 'yes', 'tanh', "adam", 'yes', 98.22])
x.add_row([3, 'no', 'relu', 'adam', 'no', 98.2])
x.add_row([3, 'yes', 'relu', 'adam', 'yes', 98.35])
x.add_row([5, 'no', 'relu', 'adam', 'no', 97.6])
x.add_row([5, 'yes', 'relu', 'adam', 'yes', 98.16])
print(x)

```



Hiddden layers	batch normalization	activation function	optimizer	dropout
2	no	relu	adam	no
2	yes	relu	adam	yes
2	yes	sigmoid	SGD	yes
2	yes	tanh	SGD	yes
2	yes	sigmoid	adam	yes
2	yes	tanh	adam	yes
3	no	relu	adam	no
3	yes	relu	adam	yes
5	no	relu	adam	no
5	yes	relu	adam	yes

Conclusion: 1.We can see that 2 hidden layer with batch normalization and dropout gives higher accuracy. 2.If the models tend to overfit. 3.Without doing batch normalization the convergence rate is slow,so batch normalization without dropout and only batch normalization we can see it performs better than a MLP which doesn't have both batch normalization and dropout. 4. Adam as best optimizer as compared to SGD. 5. Tanh activation function with ! with SGD optimizer. 6. ReLU is best activation function. 7. So for MNIST dataset, a MLP works well when it has dropout and 2 hidden layers

This file was updated remotely or in another tab. To force a save, overwriting the last update, select Save from the File menu

